

STT 461 Homework 3 Solutions

Paul Speaker

2023-03-12

1. Define a probability density function as

$$p(x) = cx^3$$

on the interval $0 < X < 3$.

- Determine what c is.

$$\int_0^3 x^3 dx = 81/4$$

So $c = 4/81$

- Use the inverse cdf method to simulate 10,000 outcomes of the distribution. Estimate the mean and standard deviation of the distribution.

$$cdf = \frac{1}{81}x^4$$

$$q(x) = \sqrt[4]{81x}$$

```
samp <- (81*runif(10000))^0.25  
mean(samp)
```

```
## [1] 2.407197
```


```
sd(samp)
```

```
## [1] 0.4861339
```

- Plot a histogram of your simulation results. Does it match what you would expect?


```
library(ggplot2)  
sampdf <- data.frame("samp"=samp)  
ggplot(data = sampdf, aes(x = samp)) + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



d. Re-do (c) with geom_density instead of geom_histogram. What is it doing?

```
ggplot(data = sampdf, aes(x = samp)) + geom_density()
```



This is a smoothed histogram. This is also an approximation for kernel density.

- e. Numerically approximate the expected value, variance, and standard deviation of the distribution.
Calculate the expected value from the definition with an integral and compare how close the 2 are.
(This duplicates a previous part, but I'll do it with more samples)

```
samp <- (81*runif(1000000))^0.25
mean(samp)
```

```
## [1] 2.399445
```

```
var(samp)
```

```
## [1] 0.2399237
```

```
sd(samp)
```

```
## [1] 0.48982
```

From the integral

$$\int_0^3 xp(x)dx = \int_0^3 \frac{4}{81}x^4 = \frac{4}{81} \cdot \frac{243}{5} = 2.4$$

- f. Use a simulation with acceptance/rejection to approximate the expectation of X, given that x is greater than 2.

There are many ways to throw out the smaller values. This is a shortcut without a loop:

```
samp_lim <- ifelse(samp>2,samp, NA)
mean(na.omit(samp_lim))
```


```
## [1] 2.596488
```

2. For the following integral :

$$\int_0^8 \frac{e^{\sin x}}{1 + \sqrt{x}} dx$$

a. Generate a plot of the region.

```
x <- seq(0,8, by = 0.01)
y <- exp(sin(x))/(1+sqrt(x))
plot_df <- data.frame("x"=x, "y"=y)
ggplot(data = plot_df, aes(x=x, y=y)) + geom_line()
```



b. Estimate using standard Monte Carlo integration ($n = 100,000$)

```
x <- runif(100000,0,8)
y <- exp(sin(x))/(1+sqrt(x))
integral1 <- 8*mean(y)
integral1
```

```
## [1] 4.423017
```

c. Estimate using acceptance/rejection Monte Carlo (again, $n = 100,000$)

```
x <- runif(100000,0,8)
y <- runif(100000,0,1.25)
prop <- sum(y<exp(sin(x))/(1+sqrt(x)))/100000
integral2 <- prop*8*1.25
integral2
```

```
## [1] 4.3911
```

3. For the following integral:

$$\int_{-\infty}^{\infty} e^{-(x-6)^4}$$

- a. Compute the following integral both with the standard Monte Carlo approach (with appropriately truncated bounds) and with Importance Sampling, using an appropriate normal distribution as an importance function. Use n = 10,000 both times.

```
# Standard Monte Carlo
x <- runif(10000,0, 12)
y <- exp(-1*(x-6)^4)
integral1 <- 12*mean(y)
integral1

## [1] 1.800777

# Importance Sampling. Using Gaussian centered at 6 for importance function.
vals <- rnorm(10000,6,1)
integral2 <- mean(exp(-1*(vals-6)^4)/dnorm(vals,6,1))
integral2

## [1] 1.823348
```

- b. Re-do both methods 100 times. Compute the variance of the answers from each method and compare.

```
# Standard Monte Carlo
integral1 <- rep(0,100)
for(i in 1:100){
  x <- runif(10000,0, 12)
  y <- exp(-1*(x-6)^4)
  integral1[i] <- 12*mean(y)
}
mean(integral1)

## [1] 1.811642
var(integral1)

## [1] 0.001216687

# Importance Sampling. Using Gaussian centered at 6 for importance function.
integral1 <- rep(0,100)
for(i in 1:100){
  vals <- rnorm(10000,6,1)
  integral2[i] <- mean(exp(-1*(vals-6)^4)/dnorm(vals,6,1))
}
mean(integral2)

## [1] 1.812771
var(integral2)

## [1] 9.815365e-05
```

We can see that the result with importance sampling has about 1/10 the variance, so it is a more accurate method.

4. Suppose we have an exponential probability distribution with parameter lambda = 1/5, so that its expected value is 5. We will perform some simulation experiments to determine the behavior of the sample mean.

- a. Generate 2,000 simulations which each constitute 100 outcomes of this distribution. Calculate the mean values for the 2,000 experiments.
- b. Next, calculate the mean value and variance for the 2,000 sample means.

```
#for (a) and (b)
sims <- 2000
sampszie <- 100
mn <- replicate(sims,0)
for(i in 1:sims){
  mn[i] <- mean(rexp(sampszie, 1/5))
}
mn_100 <- mean(mn)
var_100 <- var(mn)
mn_100

## [1] 5.003865
var_100

## [1] 0.2533212
```

- c. Repeat (a) and (b) for $N = 1,000, 10,000, 100,000$, and $1,000,000$ outcomes (2,000 simulations of each). Note: 1,000,000 could take several minutes.
- | N | Mean of Sample Means | Variance of Sample Means |
|-----------|----------------------|--------------------------|
| 1,000 | 100 | 100 |
| 10,000 | 100 | 100 |
| 100,000 | 100 | 100 |
| 1,000,000 | 100 | 100 |

```
sims <- 2000
sampszie <- 1000
mn <- replicate(sims,0)
for(i in 1:sims){
  mn[i] <- mean(rexp(sampszie, 1/5))
}
mn_1000 <- mean(mn)
var_1000 <- var(mn)
sampszie <- 10000
mn <- replicate(sims,0)
for(i in 1:sims){
  mn[i] <- mean(rexp(sampszie, 1/5))
}
mn_10000 <- mean(mn)
var_10000 <- var(mn)

sampszie <- 100000
mn <- replicate(sims,0)
for(i in 1:sims){
  mn[i] <- mean(rexp(sampszie, 1/5))
}

mn_100000 <- mean(mn)
var_100000 <- var(mn)
sampszie <- 1000000
mn <- replicate(sims,0)
for(i in 1:sims){
  mn[i] <- mean(rexp(sampszie, 1/5))
}
```

```

mn_1000000 <- mean(mn)
var_1000000 <- var(mn)

meanmean <- c(mn_100, mn_1000, mn_10000, mn_100000, mn_1000000)
varmean <- c(var_100, var_1000, var_10000, var_100000, var_1000000)
meanmean

## [1] 5.003865 5.001766 4.998906 4.999732 4.999997


```

- (d) Plot the $\log(\text{Variance})$ vs. $\log(N)$. Do you get something which is close to a straight line? What is the slope? Does the relationship between N and the variance close to what you would expect with the Central Limit Theorem?

```

samp_df <- data.frame("n" = 10^(2:6), "variance" = varmean)
plot(log(samp_df$n), log(samp_df$variance))

```



```
summary(lm(data = samp_df, log(variance) ~ log(n)))
```

```

##
## Call:
## lm(formula = log(variance) ~ log(n), data = samp_df)
##
## Residuals:
##      1       2       3       4       5 
## -0.004117  0.009588  0.002006 -0.016310  0.008833 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 5.00000   0.00000  5.0000 0.0000000 *** 
## log(n)      -0.004117  0.000000 -4.1170 0.0000000 *** 
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
```

```
## (Intercept) 3.272257  0.016558   197.6 2.86e-07 ***
## log(n)      -1.007832  0.001695  -594.6 1.05e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01234 on 3 degrees of freedom
## Multiple R-squared:      1, Adjusted R-squared:      1
## F-statistic: 3.536e+05 on 1 and 3 DF,  p-value: 1.049e-08
```

Note: lm is not necessary, just to show relationship Since the slope is -1, the variance experimentally is inversely proportional to n, as we would expect with the Central Limit Theorem.