

STT481 HW4

100 points total


Q1. Load the Boston dataset

```
set.seed(1)
library(MASS)
attach(Boston)
```

(a)

```
lm.fit = lm(nox ~ poly(dis, 3), data = Boston)

dislim = range(dis)
dis.grid = seq(from = dislim[1], to = dislim[2], by = 0.1)
lm.pred = predict(lm.fit, list(dis = dis.grid))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.grid, lm.pred, col = "red", lwd = 2)
```




(b) When the degree is 4, it results in the minimum test error.

```
library(boot)
all.deltas = rep(NA, 10)
for (i in 1:10) {
```

```

  glm.fit = glm(nox ~ poly(dis, i), data = Boston)
  all.deltas[i] = cv.glm(Boston, glm.fit, K = 10)$delta[1]
}
plot(1:10, all.deltas, xlab = "Degree", ylab = "CV error", type = "l", pch = 20, lwd = 2)

```




The polynomial regression with the best degree is given below.

```

lm.fit = lm(nox ~ poly(dis, which.min(all.deltas)), data = Boston)
lm.pred = predict(lm.fit, list(dis = dis.grid))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.grid, lm.pred, col = "red", lwd = 2)


```



(c) We fit a cubic spline with 6 degrees of freedom.

```
library(splines)
sp.fit = lm(nox ~ bs(dis, df = 6), data = Boston)

sp.pred = predict(sp.fit, list(dis = dis.grid))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.grid, sp.pred, col = "red", lwd = 2)
```



```
attr(bs(dis, df = 6), "knots")
```


```
##      25%      50%      75%
## 2.100175 3.207450 5.188425
```

The knot locations are 2.100175, 3.207450, and 5.188425.

(d) We use a 10-fold cross-validation to find best df. We try all integer values of df between 4 and 16.

```
set.seed(1)
all.df <- 4:16
all.cv <- rep(0, length(all.df))
counter <- 0
for (df in all.df) {
  counter <- counter + 1
  lm.fit <- glm(nox ~ bs(dis, df = df), data = Boston)
  all.cv[counter] <- cv.glm(Boston, lm.fit, K = 10)$delta[1]
}

plot(all.df, all.cv, lwd = 2, type = "l", xlab = "df", ylab = "CV error")
```



```
print(all.df[which.min(all.cv)])
```

```
## [1] 7
```


CV error is more jumpy in this case, but attains minimum at df=7. We pick 7 as the optimal degrees of freedom.

The cubic spline with the best degrees of freedom is given below.

```

lm.fit <- glm(nox ~ bs(dis, df = all.df[which.min(all.cv)]), data = Boston)
sp.pred <- predict(sp.fit, list(dis = dis.grid))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.grid, sp.pred, col = "red", lwd = 2)

```




(e) We fit a natural cubic spline with 6 degrees of freedom.

```

library(splines)
sp.fit = lm(nox ~ ns(dis, df = 6), data = Boston)

sp.pred = predict(sp.fit, list(dis = dis.grid))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.grid, sp.pred, col = "red", lwd = 2)

```



```
attr(ns(dis, df = 6), "knots")
```


```
## 16.66667% 33.33333%      50% 66.66667% 83.33333%
## 1.851317 2.384033 3.207450 4.325700 6.062200
```

The knot locations are 1.851317, 2.384033, 3.207450, 4.325700, and 6.062200.

(f) We use a 10-fold cross-validation to find best df. We try all integer values of df between 4 and 16.

```
set.seed(1)
all.df <- 4:16
all.cv <- rep(0, length(all.df))
counter <- 0
for (df in all.df) {
  counter <- counter + 1
  lm.fit <- glm(nox ~ ns(dis, df = df), data = Boston)
  all.cv[counter] <- cv.glm(Boston, lm.fit, K = 10)$delta[1]
}

plot(all.df, all.cv, lwd = 2, type = "l", xlab = "df", ylab = "CV error")
```




```
print(all.df[which.min(all.cv)])
```

```
## [1] 9
```

CV error is more jumpy in this case, but attains minimum at $df=9$. We pick 9 as the optimal degrees of freedom.

The natural cubic spline with the best degrees of freedom is given below.

```
lm.fit <- glm(nox ~ ns(dis, df = all.df[which.min(all.cv)]), data = Boston)
sp.pred <- predict(sp.fit, list(dis = dis.grid))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.grid, sp.pred, col = "red", lwd = 2)
```



Q2

```
library(ISLR)
data("College")
set.seed(123)
train <- sample(nrow(College), 600)
College.train <- College[train, ]
College.test <- College[-train, ]
```

(a)

```
library(leaps)
reg.fit = regsubsets(Outstate ~ ., data = College.train, nvmax = 17, method = "forward")
reg.summary = summary(reg.fit)
par(mfrow = c(1, 3))
min.cp = which.min(reg.summary$cp)
print(min.cp)
```

```
## [1] 12
```

```
plot(reg.summary$cp, xlab = "Number of Variables", ylab = "Cp", type = "l")
points(min.cp, reg.summary$cp[min.cp], pch = 4, col = "red", lwd = 7)

min.bic = which.min(reg.summary$bic)
print(min.bic)
```


```
## [1] 10
```

```
plot(reg.summary$bic, xlab = "Number of Variables", ylab = "BIC", type = "l")
points(min.bic, reg.summary$bic[min.bic], pch = 4, col = "red", lwd = 7)
```

```
max.adjr2 = which.max(reg.summary$adjr2)
print(max.adjr2)
```

```
## [1] 14
```

```
plot(reg.summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted R2",
     type = "l", ylim = c(0.4, 0.84))
points(max.adjr2, reg.summary$adjr2[max.adjr2], pch = 4, col = "red", lwd = 7)
```



```
coefi = coef(reg.fit, id = min.bic)
names(coefi)
```

```
## [1] "(Intercept)" "PrivateYes"   "Apps"          "Accept"        "Enroll"
## [6] "Top10perc"    "Room.Board"   "PhD"           "perc.alumni"   "Expend"
## [11] "Grad.Rate"
```

Cp and Adj R2 show that size 12 and 14 are the minimum sizes for the subset, respectively, while BIC chooses the size = 10. We pick 10 as the best subset size based on BIC criterion but you can also use Cp or Adj R2 as your criterions.

(b)

```


library(gam)

## Loading required package: foreach

## Loaded gam 1.20.2

gam.fit = gam(Outstate ~ Private + s(Apps, df = 5) + s(Accept, df = 5) + s(Enroll, df = 5) +
              s(Top10perc, df = 5) + s(Room.Board, df = 5) + s(Personal, df = 5) +
              s(perc.alumni, df = 5) + s(Expend, df = 5) + s(Grad.Rate, df = 5), data = College.train)
par(mfrow = c(3, 4))
plot(gam.fit, se = T, col = "red")

```



Findings: From the plots, it appears that higher enroll numbers result in lower out-of-state tuition fee. It also appears that “Accept”, “Enroll”, “Room.Board”, and “perc.alumni” have a linear relationship with out-of-state tuition. “Private” predictor indicates private and public universities have significantly different out-of-state tuition by holding other predictors. For the predictors “Apps”, “Top10perc”, “Personal”, “Expend” and “Grad.Rate”, it seems there exists a quadratic or higher degrees effect. Note here I chose $\text{df} = 5$ but you can use other dfs. In general, $\text{df} = 3-6$ would be recommended.

(c)

```

gam.pred = predict(gam.fit, College.test)
gam.err = mean((College.test$Outstate - gam.pred)^2)
gam.err

```

```

## [1] 3183501

```

We obtain a test mean squared error of 3183501 using GAM with 10 predictors.

We can use `summary` to see whether the non-linear effects are significant or not.


```
summary(gam.fit)

##
## Call: gam(formula = Outstate ~ Private + s(Apps, df = 5) + s(Accept,
##       df = 5) + s(Enroll, df = 5) + s(Top10perc, df = 5) + s(Room.Board,
##       df = 5) + s(Personal, df = 5) + s(perc.alumni, df = 5) +
##       s(Expend, df = 5) + s(Grad.Rate, df = 5), data = College.train)
## Deviance Residuals:
##      Min      1Q   Median      3Q     Max
## -6275.98 -1069.78    85.95  1147.84  7446.96
##
## (Dispersion Parameter for gaussian family taken to be 3210667)
##
## Null Deviance: 9694792870 on 599 degrees of freedom
## Residual Deviance: 1775499257 on 553.0002 degrees of freedom
## AIC: 10738.98
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##                               Df   Sum Sq   Mean Sq F value    Pr(>F)
## Private                      1 2742410325 2742410325 854.156 < 2.2e-16 ***
## s(Apps, df = 5)                1 1195413597 1195413597 372.326 < 2.2e-16 ***
## s(Accept, df = 5)               1  97976877  97976877 30.516 5.103e-08 ***
## s(Enroll, df = 5)                1 171618769 171618769 53.453 9.337e-13 ***
## s(Top10perc, df = 5)             1  815049317  815049317 253.857 < 2.2e-16 ***
## s(Room.Board, df = 5)              1 491831182 491831182 153.187 < 2.2e-16 ***
## s(Personal, df = 5)                1  39418544  39418544 12.277 0.0004956 ***
## s(perc.alumni, df = 5)              1 219360932 219360932 68.323 1.035e-15 ***
## s(Expend, df = 5)                  1  574736351  574736351 179.008 < 2.2e-16 ***
## s(Grad.Rate, df = 5)                 1  43959989  43959989 13.692 0.0002369 ***
## Residuals                     553 1775499257   3210667
##
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##                               Npar Df   Npar F    Pr(F)
## (Intercept)
## Private
## s(Apps, df = 5)                  4  2.1759 0.0704230 .
## s(Accept, df = 5)                 4 21.4399 2.22e-16 ***
## s(Enroll, df = 5)                 4  5.2582 0.0003665 ***
## s(Top10perc, df = 5)               4  1.2001 0.3096976
## s(Room.Board, df = 5)               4  1.9962 0.0937235 .
## s(Personal, df = 5)                 4  2.8825 0.0221169 *
## s(perc.alumni, df = 5)               4  1.0929 0.3591952
## s(Expend, df = 5)                  4 25.3298 < 2.2e-16 ***
## s(Grad.Rate, df = 5)                 4  3.3488 0.0100772 *
##
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

To avoid overfitting issue, you can also remove the insignificant nonlinear effects, such as `perc.alumni`, where the p-value is 0.3591952. For example, you can do so by replacing `s(perc.alumni, df=5)` with `perc.alumni`. See below.

```
gam.fit = gam(Outstate ~ Private + Apps + s(Accept, df = 5) + s(Enroll, df = 5) +
              Top10perc + Room.Board + s(Personal, df = 5) + perc.alumni +
              s(Expend, df = 5) + s(Grad.Rate, df = 5), data = College.train)
# I leave Grad.Rate as it is since it's not too insignificant
par(mfrow = c(3, 4))
plot(gam.fit, se = T, col = "red")
gam.pred = predict(gam.fit, College.test)
gam.err = mean((College.test$Outstate - gam.pred)^2)
gam.err
```

`## [1] 2846868`



We obtain a lower test mean squared error of 2846868 by removing the insignificant nonlinear effects.

Q3 (a)

```
train.dat <- read.csv("nbadata17.csv", row.names = 1)
test.dat <- read.csv("nbadata18.csv", row.names = 1)
```


(b)

```
library(leaps)
regfit.fwd <- regsubsets(salary ~ ., data = train.dat, nvmax = 32, method = "forward")
```

```

reg.fwd.summary <- summary(regfit.fwd)
plot(reg.fwd.summary$bic, xlab = "Number of Variables ", ylab = "BIC", type = "l")
points(which.min(reg.fwd.summary$bic), reg.fwd.summary$bic[which.min(reg.fwd.summary$bic)], col = "red",

```



```
coef(regfit.fwd, 5)
```

```

##   (Intercept)      age       g      gs  x2ppercnt      pts
## -10003.77541  360.81836 -69.17732  79.48794  8599.43163  644.44076

```

BIC criterion indicates it selects 5 predictors, and they are `age`, `g`, `gs`, `x2ppercnt`, and `pts`.

(c) The prediction of their salaries are below.

```

library(leaps)
predict.regsubsets <- function (object, newdata , id, ...){
  form <- as.formula(object$call[[2]]) # formula of null model
  mat <- model.matrix(form, newdata) # building an "X" matrix from newdata
  coefi <- coef(object, id = id)      # coefficient estimates associated with the object model contain
  xvars <- names(coefi)               # names of the non-zero coefficient estimates
  return(mat[,xvars] %*% coefi)     # X[,non-zero variables] %*% Coefficients[non-zero variables]
}
fwd.pred <- predict.regsubsets(regfit.fwd, newdata = test.dat, id = which.min(reg.fwd.summary$bic))
fwd.pred

##          [,1]
## Khris Middleton 16496.228
## Terrence Ross    8028.138
## Ish Smith        6595.110
## Myles Turner     11748.581
## Thaddeus Young   14625.969

```


(d) We fit a ridge regression model on the training set and predict the salaries using the fitted model.

```
library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-4

X <- model.matrix(salary ~ . , train.dat)[,-1] # the first column is the intercept
y <- train.dat$salary
ridge.mod <- glmnet(X, y, alpha = 0) # 0 indicates ridge regression
cv.out <- cv.glmnet(X, y, alpha = 0, nfolds = 10)
plot(cv.out)
```



```
bestlam <- cv.out$lambda.min
bestlam


## [1] 1699.232

test.X <- model.matrix(salary ~ ., test.dat) # same form as X
ridge.pred <- predict(ridge.mod, s = bestlam, newx = test.X)
ridge.pred
```

```
##          s1
## Khris Middleton 16369.088
## Terrence Ross    9126.686
## Ish Smith       6365.419
## Myles Turner     12378.610
## Thaddeus Young   12827.164
```

- (e) We fit a lasso model on the training set and predict the salaries using the fitted model, and report the non-zero coefficient estimates.

```
lasso.mod <- glmnet(X, y, alpha = 1)
cv.out <- cv.glmnet(X, y, alpha = 1)
plot(cv.out)
```



```
bestlam <- cv.out$lambda.min
coef(lasso.mod, s = bestlam)
```

```
## 31 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) -5887.8085261
## posPF       .
## posPG      -955.1017882
## posSF      129.0623908
## posSG     -172.5341187
## age        316.9460181
## g         -48.6773725
## gs        69.4018629
## mp        .
## fg        367.0638893
## fga       .
## fgpercent -1716.0581083
## x3p       .
## x3pa      684.2299547
## x3ppercnt -154.8122753
## x2p       0.4944476
## x2pa      .
```

```

## x2ppercen 9281.5959872
## efgpercent .
## ft 1747.9818535
## fta .
## ftpercent -4429.0627815
## orb 564.5035918
## drb .
## trb 252.6887510
## ast 1064.3383195
## stl .
## blk 982.4325082
## tov -1933.3153465
## pf -703.1011589
## pts .
lasso.pred <- predict(lasso.mod, s = bestlam, newx = test.X)
lasso.pred

```

```

##          s1
## Khris Middleton 16471.288
## Terrence Ross    8022.151
## Ish Smith      5884.402
## Myles Turner    12878.233
## Thaddeus Young 13362.675

```


- (f) We fit a GAM on the training set using the five predictors: `age`, `g`, `gs`, `x2ppercen`, and `pts`, and predict the salaries using the fitted model.

```

library(gam)
gam.fit <- gam(salary ~ s(age, df = 5) + s(g, df = 5) + s(gs, df = 5) +
                 s(x2ppercen, df = 5) + s(pts, df = 5),
                 data = train.dat)
par(mfrow = c(2, 3))
plot(gam.fit, se = T, col = "red")
predict(gam.fit, test.dat)

```

<code>## Khris Middleton</code>	<code>Terrence Ross</code>	<code>Ish Smith</code>	<code>Myles Turner</code>	<code>Thaddeus Young</code>
<code>## 18961.822</code>	<code>8257.568</code>	<code>7890.952</code>	<code>9736.555</code>	<code>15201.548</code>



Findings: The `age` appears to have a quadratic effect: if the player is young or old, the salary is low, and if the player is at middle age, the salary is high. The `gs` have a linear effect: if the player plays more games as a starter, then salary is high. Similarly `g` and `x2ppercen` appears to have non-linear effects, and `pts` has a linear effect (which is not too surprising: if a player scores more, then the salary is higher).

Similarly, we can use `summary` to see whether the non-linear effects are significant or not.

```
summary(gam.fit)
```

```
##
## Call: gam(formula = salary ~ s(age, df = 5) + s(g, df = 5) + s(gs,
##      df = 5) + s(x2ppercen, df = 5) + s(pts, df = 5), data = train.dat)
## Deviance Residuals:
##    Min      1Q Median      3Q     Max
## -17058   -3134   -643   2841   17099
##
## (Dispersion Parameter for gaussian family taken to be 32577645)
##
## Null Deviance: 23146145548 on 340 degrees of freedom
## Residual Deviance: 10261941621 on 314.9995 degrees of freedom
## AIC: 6893.678
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##              Df     Sum Sq   Mean Sq F value    Pr(>F)
## s(age, df = 5)  1 1.5450e+09 1544968977 47.4242 3.109e-11 ***
## s(g, df = 5)    1 1.2687e+09 1268653996 38.9425 1.407e-09 ***
## s(gs, df = 5)   1 5.8595e+09 5859484135 179.8621 < 2.2e-16 ***
## s(x2ppercen, df = 5) 1 1.4296e+08 142959048   4.3883  0.03699 *
```

```

## s(pts, df = 5)           1 2.2529e+09 2252932342 69.1558 2.772e-15 ***
## Residuals                 315 1.0262e+10   32577645
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##                               Npar Df Npar F      Pr(F)
## (Intercept)                   4    6.8175 2.837e-05 ***
## s(age, df = 5)                4  1.5708  0.181864
## s(gs, df = 5)                 4  0.8509  0.493884
## s(x2ppercents, df = 5)       4  4.0478  0.003244 **
## s(pts, df = 5)                 4  3.5543  0.007456 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

It appears that `g` and `gs`'s non-linear effects are very insignificant. If we remove the non-linear effects of `g` and `gs`, then the prediction is below.

```

gam.fit <- gam(salary ~ s(age, df = 5) + g + gs +
                 s(x2ppercents, df = 5) + s(pts, df = 5),
                 data = train.dat)
gam.pred <- predict(gam.fit, test.dat)
gam.pred

## Khris Middleton   Terrence Ross       Ish Smith     Myles Turner Thaddeus Young
##          18319.679      8471.934      7070.722      9352.009      14714.949

```